


CORE-Bench: Comprehensive LLM Reasoning Evaluation Dashboard



KEY FINDINGS FROM CORE-BENCH EVALUATION

- OVERALL PERFORMANCE
 - 4 of 5 models achieved 96.08% accuracy
 - Claude Haiku 4.5 scored 94.12% (lowest)
 - All models passed core reasoning tasks

- BENCHMARK STATISTICS
 - Total: 91 reasoning problems
 - 4 categories tested
 - 6 evaluation tasks + 1 comprehensive

- TOP PERFORMERS
 - Gemini 3 Flash: Perfect on all 6 tasks
 - Qwen 3 Next 80B: Perfect on all 6 tasks
 - Both demonstrated robust reasoning

- ADVANCED PROBLEMS
 - 40 advanced problems added (44% of total)
 - Simpson's Paradox, Modal Logic, etc.
 - Designed to challenge frontier models

- CHALLENGE AREAS
 - Quality Evaluation: 2 failures (40%)
 - Multi-Step Planning: 1 failure (20%)
 - Core reasoning tasks: 0 failures

- RECOMMENDATIONS
 - Re-run with expanded 91-problem set
 - Focus testing on Quality Evaluation
 - Add per-problem accuracy tracking